

Beyond Word Embeddings: Representations for Multi-modal Data

Luis Armona*, José González Brenes†, and Ralph Edezath‡

Chegg Engineering
3990 Freedom Cir, Santa Clara, CA 95054

Abstract

Methods that calculate dense vector representations for text have proven to be very successful for knowledge representation. We study how to estimate dense representations for multi-modal data (e.g., text, continuous, categorical). We propose Feat2Vec as a novel model that supports supervised learning when explicit labels are available, and self-supervised learning when there are no labels. Feat2Vec calculates embeddings for data with multiple feature types enforcing that all embeddings exist in a common space. We believe that we are the first to propose a method for learning self-supervised embeddings that leverage the structure of multiple feature types. Our experiments suggest that Feat2Vec outperforms previously published methods, and that it may be useful for avoiding the cold-start problem.

1 Introduction

Informally, in machine learning a *dense representation*, or *embedding* of a vector $\vec{x} \in \mathbb{R}^n$ is another vector $\vec{\beta} \in \mathbb{R}^r$ that has much lower dimensionality ($r \ll n$) than the original representation. In general, we consider two kind of models that produce embeddings: **(i) supervised methods**, like matrix factorization, calculate embeddings that are highly tuned to a prediction task. For example, in the Netflix challenge, movie identifiers are embedded to predict user ratings. On the other hand, **(ii) self-supervised methods** (sometimes referred to as unsupervised methods) are not tuned for the prediction task they are ultimately used for. For example, word embedding algorithms such as Word2Vec (Mikolov et al., 2013b) are self-supervised. These algorithms are typically evaluated by analogy solving, or sentiment analysis (Le and Mikolov, 2014), even though their loss functions are not tuned for either of these tasks.

In this paper we propose Feat2Vec as a novel method that embeds arbitrary feature types, such as text, numer-

ical or categorical data, in a common vector space—for both supervised and self-supervised scenarios. We believe that this is the first self-supervised algorithm that is able to calculate embeddings from data with multiple feature types. Consider the non-trivial work that was required to extend a model like Word2Vec to support additional features. The authors of the seminal Doc2Vec (Le and Mikolov, 2014) paper needed to design both a new neural network and a new sampling strategy to add a single feature (document ID). Feat2Vec is a general method that allows calculating embeddings for any number of features.

2 Feat2Vec

In this section we describe how Feat2Vec learns embeddings of feature groups.

2.1 Model

Feat2Vec predicts a target output \hat{y} from a list of feature groups constructed from a partition $\vec{\kappa}$ of raw features \vec{x} :

$$\vec{x} = \langle \vec{x}_{\vec{\kappa}_1}, \vec{x}_{\vec{\kappa}_2}, \dots, \vec{x}_{\vec{\kappa}_n} \rangle = \langle \vec{g}_1, \vec{g}_2, \dots, \vec{g}_n \rangle \quad (1)$$

Each of the n feature groups \vec{g}_i is defined a priori by $\vec{\kappa}_i$, which indexes the dimensions in \vec{x} that belong to the i -th group. The number of dimensions of each group may vary, but all of the feature groups are embedded to the same space via their *feature extraction function* ϕ_i . A feature extraction function ϕ acts on the features in κ_i , and the other features interact with the words only via the output of ϕ . These functions learn how to embed each vector (feature group) and outputs a vector in \mathbb{R}^r . More, formally:

$$\hat{y}(\vec{x}, \vec{\phi}, \vec{\kappa}) = \omega \left(\sum_{i=1}^n \sum_{j=i}^n \overbrace{\phi_i(\vec{x}_{\vec{\kappa}_i})}^{\text{group } i \text{ embedding}} \cdot \overbrace{\phi_j(\vec{x}_{\vec{\kappa}_j})}^{\text{group } j \text{ embedding}} \right) \quad (2)$$

Here, ω is an activation function. Intuitively, the dot product (\cdot) returns a scalar that measures the (dis)similarity between the embeddings of the feature groups. A simple implementation of ϕ_i is a linear fully-connected layer, where the output of the r -th entry is:

*Email: larmona@stanford.edu. Affiliation: Department of Economics, Stanford University

†Email: jgonzalez@chegg.com

‡Email: redezath@chegg.com

Copyright © 2015, Association for the Advancement of Artificial Intelligence (www.aaai.org). All rights reserved.

$$\phi_{i,r}(\vec{x}_i; \vec{\theta}_r)_r = \sum_{a=1}^{d_i} \theta_{r_a} x_{i_a} \quad (3)$$

Other functional forms of ϕ_i we have explored in our experiments include:

- A convolutional neural network that transforms a text sequence into a r -dimensional embedding. This function is able to preserve information in the order of a text sequence, unlike other methods for transforming text to numerical values such as bag-of-words.
- A deep, fully connected neural network that projects a scalar, such as average rating, into a r -dimensional embedding space. This extraction function, by treating the input as numerical, will require scalars close in value to have similar embeddings.

Any neural function that outputs a r -dimensional vector could be plugged into our framework as an extraction function ϕ .

Figure 1 compares existing factorization methods with our novel model. Figure 1a shows a Factorization Machine (Rendle, 2010) that embeds each feature. Figure 1b shows Feat2Vec with two feature groups: the first group only has a single feature which is projected to an embedding (just like matrix factorization); but the second group has multiple features, which are together projected to a single embedding. Figure 1c shows an approach of using neural networks within factorization machines that has been proposed multiple times (Dziugaite and Roy, 2015; Guo et al., 2017). It replaces the dot product of factors with a learned neural function, which has been shown to improve predictive accuracy. The caveat of this architecture is that is no longer possible to interpret the embeddings as latent factors related to the target task.

The structure of the Feat2Vec model significantly expands previously published models. In traditional factorization methods, item identifiers are embedded and thus need to be observed during training (i.e., cold-start problem). Feat2Vec can learn an embedding from an alternative characterization of an item, such as a textual description. Feat2Vec extends Factorization Machine (Rendle, 2010), in that we allow calculating embeddings for feature groups and we use feature extraction functions. StarSpace (Wu et al., 2018) introduces feature groups (called entities in their paper), but constrains them to be “bags of features”, is supervised, and limited to two feature types. In contrast, Feat2Vec allows continuous features, and does not require the user to assign one feature as a specific target. Additionally, we propose a self-supervised learning algorithm.

2.2 Supervised Learning from Data

We can learn the parameters of a Feat2Vec model θ using training data by minimizing a loss function \mathcal{L} :

$$\arg \min_{\theta} \sum_x \mathcal{L}(y(x), \hat{y}(x; \theta)) + \gamma \|\theta\|^w \quad (4)$$

Here, $y(x)$ is the true target value for x obtained from training data, and $\hat{y}(x)$ is the one estimated by the model (Equation 2); θ represents the parameters learned in during training (i.e. the parameters associated with the extraction functions ϕ_i). The hyperparameter γ controls the amount of regularization. For labeling and classification tasks, we optimize the binary cross-entropy.

It is straightforward to optimize Equation 4 directly. Typically, this loss will be for a prediction task, and involve at least two feature groups—one of the feature groups is the target label that we want to predict, such as rating, and the other group(s) is the input from which we want to make the prediction, such as review text.

Consider a target feature group g_i that exists in a high dimensional space. This could be a one-hot encoding of a categorical variable with a large number of possible values. In such a scenario, it is overwhelmingly costly to feed the model all negative labels, particularly if the model is sparse. In these cases, we may instead use a binary classifier with implicit sampling (Dyer, 2014). In implicit sampling, instead of using all of the possible negative labels, one simply samples a fixed number (k) from the set of possible negative labels for each positively labelled record.

2.3 Self-supervised Learning From Data

We now discuss how Feat2Vec can be used to learn embeddings in an self-supervised setting with no explicit target for prediction.

The training dataset for a Feat2Vec model consists of the observed data. In natural language, these would be documents written by humans, along with document metadata. Since self-supervised Feat2Vec will require positive and negative examples during training, we supply unobserved data as negative examples. For example, Word2Vec samples a negative observation from a noise distribution \mathcal{Q}_{w2v} that is proportional to the empirical frequency of a word in the training data.

We introduce a new implicit sampling method that enables learning self-supervised embeddings for feature groups. Unlike Word2Vec, we do not constraint features types to be words. Instead, by grouping features using the parameter $\vec{\kappa}$ in Equation 2, the model can reason on more abstract entities in the data. By entity, we mean a particular feature group value. For example, in our experiments on a movie dataset, we define a “genre” feature group, where we group non-mutually exclusive indicators for movie genres, including comedy, action, and drama films.

We start with a dataset S^+ of records with n feature groups. We then mark all observed records in the training set as our positive-labeled examples. For each positive record, we generate k negative labels using the 2-step algorithm documented in Algorithm 1:

Explained in words, our negative sampling method for self-supervised learning iterates over all of the observations of the training dataset. For each observation \vec{x}^+ , it randomly selects the i -th feature group from a noise distribution $\mathcal{Q}_1(\cdot)$. Then, it creates a negative




Figure 1: **Network architectures for factorization models.** The white clouds (☁) represent deep layers, for example a convolutional network for text features, while the dots (\cdot) denote dot products.

Algorithm 1 Implicit sampling algorithm for self-supervised Feat2Vec

```

1: function FEAT2VEC_SAMPLE( $S^+, k, \alpha_1, \alpha_2$ )
2:    $S^- \leftarrow \emptyset$ 
3:   for  $\vec{x}^+ \in S^+$  do
4:     Draw a random feature group  $\kappa_i \sim$ 
        $\mathcal{Q}_1(\{\text{params}(\phi_i)\}_{i=1}^{|\kappa|}, \alpha_1)$ 
5:     for  $j \in \{1, \dots, k\}$  do
6:        $\vec{x}^- \leftarrow \vec{x}^+$   $\triangleright$  set initially to be equal to the
                           positive sample
7:       Draw a random entity  $\tilde{x} \sim \mathcal{Q}_2(X_{\kappa_i}, \alpha_2)$ 
8:        $\vec{x}^-_{\kappa_i} \leftarrow \tilde{x}$   $\triangleright$  substitute the  $i$ -th feature group
                           with the sampled one
9:      $S^- \leftarrow S^- + \{\vec{x}^-\}$ 
10:   end for
11: end for
12: return  $S^-$ 
13: end function

```

observation that is identical to \vec{x}^+ , except that its i -th feature group value is replaced by a value sampled from a noise distribution $\mathcal{Q}_2(\cdot)$. In our application, we use the same class of noise distributions (flattened multinomial) for both levels of sampling, but this need not be the case. We now describe the noise distributions that we use. Let $P_{\mathcal{Q}}(x)$ denote the probability of x under distribution \mathcal{Q} .

Sampling Feature Groups. The function `params` calculates the complexity of a feature extraction function ϕ_i . To sample a feature group, we choose a feature group κ_i from a multinomial distribution with probabilities proportional a feature’s complexity. By complexity, we mean the number of parameters we learn that are associated with a feature group’s extraction function ϕ_i . This choice places more weight on features that have more parameters and thus are going to require more training iterations to properly learn. The sampling probabilities of each feature group are:

$$P_{\mathcal{Q}_1}(\kappa_i | \text{params}(\phi_i)_{i=1}^{|\kappa|}, \alpha_1) = \frac{\text{params}(\phi_i)^{\alpha_1}}{\sum_{j=1}^{|\kappa|} \text{params}(\phi_j)^{\alpha_1}} \quad (5)$$

For categorical variables using a linear fully-connected layer, the complexity is simply proportional to the number of categories in the feature group. However, if

we have multiple intermediate layers for some feature extraction functions (e.g., convolutional layers), these parameters should also be counted towards a feature group’s complexity. The hyper-parameter $\alpha_1 \in [0, 1]$ flattens the distribution. When $\alpha_1 = 0$, the feature groups are sampled uniformly, and when $\alpha_1 = 1$, they are sampled exactly proportional to their complexity.

Sampling Feature Group Values. To sample an entity within a feature group κ_i , we use a similar strategy to Word2Vec and use the empirical distribution of values:

$$P_{\mathcal{Q}_2}(x | X_{\kappa_i}, \alpha_2) = \frac{\text{count}(x)^{\alpha_2}}{\sum_{x'_{\kappa_i} \in S^+ \text{ count}(x'_{\kappa_i})^{\alpha_2}}, \quad \alpha_2 \in [0, 1] \quad (6)$$

Here, $\text{count}(x)$ is the number of times a value x of feature group i appeared in the training dataset S^+ , and α_2 is again a flattening hyperparameter.

This method will sometimes by chance generate negatively labeled samples that *do* exist in our sample of observed records. The literature offers two solutions: in the Negative Sampling of Word2Vec, duplicate negative samples are ignored (Dyer, 2014). Instead, we account for the probability of random negative labels being identical to positively labeled data using *Noise Contrastive Estimation* (NCE) (Gutmann and Hyvärinen, 2010).

The Loss Function for Self-Supervised Learning

For our self-supervised learning of embeddings, we optimize a NCE loss function, to adjust the structural statistical model $\hat{y} = p(y = 1 | \vec{x}, \vec{\phi}, \theta)$ expressed in Equation 2, to account for the possibility of random negative labels that appear identical to positively labeled data. Since in self-supervised learning we only deal with a dichotomous label, indicating a positive or negative sample, we restrict our attention to usage of Equation 2 with ω as a logistic link function.

An additional burden of NCE is that we need to calculate a partition function $Z_{\vec{x}}$ for each unique record type \vec{x} in the data that transforms the probability \hat{y} of a positive or negative label into a well-behaved distribution that integrates to 1. This would introduce an astronomical amount of computation and greatly increase the complexity of our model. Instead, we appeal to the work of Mnih and Teh (2012), who show that in the context of language models setting the $Z_{\vec{x}} = 1$ in advance does not change the performance of the model.

The intuition is that if the underlying model has enough free parameters, it will effectively learn the probabilities itself, since systemic under/over prediction of probabilities will result in penalties on the loss function. Written explicitly, the new structural probability model is:

$$\tilde{p}(y = 1|\vec{x}, \vec{\phi}, \boldsymbol{\theta}) = \frac{\exp(s(\vec{x}|\vec{\phi}, \boldsymbol{\theta}))}{\exp(s(\vec{x}|\vec{\phi}, \boldsymbol{\theta})) + P_Q(\vec{x}|\alpha_1, \alpha_2)} \quad (7)$$

$$s(\vec{x}|\vec{\phi}, \boldsymbol{\theta}) = \sum_{i=1}^{|\vec{\kappa}|} \sum_{j=i}^{|\vec{\kappa}|} \phi_i(\vec{x}_{\vec{\kappa}_i}) \cdot \phi_j(\vec{x}_{\vec{\kappa}_j}) \quad (8)$$

$s(\cdot)$ denotes the score of a record \vec{x} given parameters/extraction functions, and $P_Q(\cdot)$ denotes the probability of a record \vec{x} being drawn from our negative sampling algorithm, conditional on the positively labeled record \vec{x}^+ the negative sample is drawn for:

$$P_Q(\vec{x}|\alpha_1, \alpha_2, X, \vec{x}^+) = P_{Q_2}(\vec{x}_{\kappa_i}|X_{\kappa_i}, \alpha_2) \times P_{Q_1}(\kappa_i | \text{params}(\phi_i)\}_{i=1}^n, \alpha_1)$$

Our loss function L optimizes $\boldsymbol{\theta}$, the parameters of the feature extraction functions $\vec{\phi}$, while accounting for the probability of negative samples.

$$L(S) = \arg \min_{\boldsymbol{\theta}} \frac{1}{|S^+|} \sum_{\vec{x}^+ \in S^+} \left(\log(\tilde{p}(y = 1|\vec{x}^+, \vec{\phi}, \boldsymbol{\theta})) + \sum_{\vec{x}^- \sim Q(\cdot|\vec{x}^+)}^k \log(\tilde{p}(y = 0|\vec{x}^-, \vec{\phi}, \boldsymbol{\theta})) \right)$$

With n feature groups, the loss function of self-supervised Feat2Vec can be shown to be a weighted average of the losses from n distinct supervised Feat2Vec models¹, one for each feature group. In other words, it optimizes n multi-label classifiers, where each classifier is optimized for a different feature group. Intuitively, in Algorithm 1, we sample a target feature group according to P_{Q_1} , and then add that group’s supervised loss to the total self-supervised loss.

3 Empirical Results

3.1 Supervised Embeddings

We now address our working hypotheses for evaluating supervised embeddings. For all our experiments we define a development set and a single test set which is 10% of the dataset, and a part of the development set is used for early stopping or validating hyper-parameters. Since we only observe positive labels during supervised training, for each positive label in the test set we sample negative labels according to the label frequency. This ensures that if a model merely predicts the labels according to their popularity, it would have an AUC of 0.5. For

¹Proof can be found in the supplemental material at http://web.stanford.edu/~larmona/supplemental_theorem.pdf

Table 1: Supervised Yelp rating prediction

	MSE	Improvement over MF
Feat2Vec	0.480	69.2 %
DeepCoNN	1.441	19.6 %
MF (Matrix Factorization)	1.561	-

the regression task, we use mean squared error (MSE) as the evaluation metric. In preliminary experiments we noticed that regularization slows down convergence with no gains in prediction accuracy, so we avoid overfitting only by using early stopping. We share most of the code for the experiments online² for reproducibility.

For our feature extraction function ϕ for text, we use a Convolutional Neural Network (CNN) that has been shown to be effective for natural language tasks (Kalchbrenner, Grefenstette, and Blunsom, 2014). Instead of tuning the hyper-parameters, we follow previously published guidelines (Zhang and Wallace, 2017).

Comparison with alternative CNN-based text factorization We now compare with a method called DeepCoNN, a deep network specifically designed for incorporating text into matrix factorization (Zheng, Noroozi, and Yu)—which reportedly, is the state of the art for predicting customer ratings when textual reviews are available. To make results comparable, the Feat2Vec experiments use the same feature extraction function used by DeepCoNN. We evaluate on the Yelp dataset³, which consists of 4.7 million reviews of restaurants. For each user-item pair, DeepCoNN concatenates the text from all reviews for that item and all reviews by that user. The concatenated text is fed into a feature extraction function followed by a factorization machine. For Feat2Vec, we build 3 feature groups: item (restaurant) identifiers, user identifiers, and review text.

Table 1 compares our methods to DeepCoNN’s published results because a public implementation is not available. Feat2Vec provides a large performance increase when comparing the reported improvement in the mean squared error (MSE) over Matrix Factorization. Our approach is more general, and we claim that it is also more efficient. Since DeepCoNN concatenates text, when the average reviews per user is \bar{n}_u and reviews per item is \bar{n}_i , each text is duplicated on average $\bar{n}_i \times \bar{n}_u$ times per training epoch. In contrast, for Feat2Vec each review is seen only once per epoch. Thus it can be 1-2 orders of magnitude more efficient for datasets where $\bar{n}_i \times \bar{n}_u$ is large.

3.2 Self-Supervised Embeddings

The prediction of similar or related words is a commonly used method for evaluating self-supervised word embeddings (Bojanowski et al., 2017; Mikolov et al.,

²<https://goo.gl/zEQBiA>

³<https://www.yelp.com/dataset/challenge>

2013a). Since our algorithm is designed for multi-modal data, we generalize this task for feature types beyond text as follows. Given a dataset with n feature groups g_1, g_2, \dots, g_n we have trained a self-supervised Feat2Vec model on, we choose a single group g_i as the “target” and a group g_j as the predictor. Given a held-out instance, we rank all possible values of the target group by cosine similarity to the source group’s embedding $\phi_j(\vec{x}_{\vec{z}_j})$ for that instance. We then retrieve the ranking of the actual target entity in the test dataset. Rankings are evaluated with according to their mean percentile rank (MPR): $MPR = \frac{1}{N} \sum_{i=1}^N R_i / (\max_i R_i)$, where R_i is the rank of the entity under our evaluation procedure for observation i . A score of 0 would indicate perfect performance (i.e. top rank every test sample given), so lower is better under this metric. We will use this method to compare the performance of Feat2Vec algorithm and compare its performance to Word2Vec’s CBOW algorithm for learning embeddings. For Word2Vec, each record is treated as a separate document. To incorporate non-text data in Word2Vec, we tokenize all feature values and treat them as words in the Word2Vec training corpus (e.g. a movie record flagged as belonging to the horror genre has the token “genre_horror” in its Word2Vec training document). In our experiments, we set $\alpha_1 = .25$, $\alpha_2 = 0.75$, and $k = 5$. Wherever possible, we use the same hyperparameters across embedding models (Word2Vec, Doc2Vec, Feat2Vec). We chose an α_1 value that would highlight Feat2Vec’s ability to prioritize learning of low-complexity feature groups which nonetheless appear in each observation. We evaluate self-supervised Feat2Vec on 3 datasets:

Movies The Internet Movie Database (IMDB) is a publicly available dataset⁴ of information related to films, television programs and video games. We limit our experiments to data on its 465,136 movies. Table 2 summarizes our feature groups. For groups with multiple categories, we use a “Bag of categories” approach. For example, we sum embeddings of individual actors such as Tom Cruise to produce a single “actor” feature group embedding associated with a title.

Education We use a dataset from an anonymized leading technology company that provides educational services. In this proprietary dataset, we have 57 million observations and 9 categorical feature types which include textbook identifier, user identifier, school identifier, along with other proprietary features. Each observation is an interaction a user had with a textbook.

Yelp We use the Yelp dataset from our supervised experiments to evaluate the efficacy of self-supervised embeddings in ratings prediction.

Table 2: IMDB Feature Groups

Feature Type Name	Type	# of feats.
Runtime (minutes)	Real-valued	1
IMDB rating (0-10)	Real-valued	1
# of rating votes	Real-valued	1
Is adult film?	Boolean	2
Movie release year	Categorical	271
Movie title	Text	165,471
Directors	Bag of categories	174,382
Genres	Bag of categories	28
Writers	Bag of categories	244,241
Principal actors	Bag of categories	1,104,280

Results After training IMDB embeddings, we use cast members associated with movies in the test set and predict the actual director of the film. We rank directors by cosine similarity of their embedding to the title’s cast member feature group embedding. For the educational dataset, we directly retrieve the most similar textbooks to the user embedding.

Table 3 presents our evaluation results. Feat2Vec outperforms CBOW in the MPR metric. Additionally, while CBOW predicts the actual director 1.26% of the times, Feat2Vec does so 2.43% of the time, a 92% improvement in Top-1 Precision over CBOW.

Table 3: Mean percentile rank

Dataset	Feat2Vec	CBOW
IMDB	19.36%	24.15%
Educational	25.2%	29.2%

3.3 Self-Supervised Feat2Vec Performance with Continuous Inputs

We now focus on how well Feat2Vec performs on a real-valued feature. We expect this task to highlight Feat2Vec’s advantage over token-based embedding learning algorithms, such as Word2Vec. Our rating embedding extraction function (a 3-layer DNN) will require embeddings of numerically similar ratings to be close, while Word2Vec will treat two numerically different ratings as completely different entities. We evaluate the prediction of test dataset ratings by choosing the rating embedding most similar⁵ to the embedding of the movie’s director, and compute the MSE of the predicted rating in the test dataset. Feat2Vec scores an MSE of 6.6, while Word2Vec (CBOW) scores 9.3.

We also use the Yelp dataset, predicting the most similar rating embedding to the review text embeddings produced by Feat2Vec, Word2Vec (CBOW), and Doc2Vec (DM). Doc2Vec is only used for the Yelp dataset because there is no clear “document” in the IMDB dataset, while in the Yelp dataset the review text is a natural candidate. For Word2Vec and Doc2Vec, the review text

⁴<http://www.imdb.com/interfaces/>

⁵As before, the metric is cosine similarity.




Figure 2: Confusion Matrices of Self-Supervised Yelp Ratings Predictions

embedding is the average of word embeddings, analogous to the context vector used for learning in these algorithms. Figure 2 reports the confusion matrices for each model from this experiment. In general, Word2Vec is poor in its predictive power, predicting 5 stars for 97% of the test sample. Though Feat2Vec and Doc2Vec yield comparable MSE (2.94 vs. 2.92, respectively), Feat2Vec outperforms Doc2Vec by a substantial margin in classification error rate (55% vs. 73%) and mean absolute error (1.13 vs. 1.31). As made evident by the figure, Doc2Vec is unable to identify low rating reviews: only 0.4% of Doc2Vec predictions are ≤ 2 stars, despite this comprising 20% of the data. In contrast, Feat2Vec is more diverse in its predictions, and better able to identify extreme reviews.

4 Conclusion

Embeddings have proven useful in a wide variety of contexts, but they are typically built from datasets with a single feature type as in the case of Word2Vec, or tuned for a single prediction task as in the case of Factorization Machine. We believe Feat2Vec is an important step towards general-purpose methods, because it decouples feature extraction from prediction for datasets with multiple feature types, it can be self-supervised, and its embeddings are easily interpretable.

In the supervised setting, Feat2Vec is able to calculate embeddings for passages of texts. We show results outperforming an algorithm specifically designed for text—even when using the same feature extraction CNN.

In the self-supervised setting, Feat2Vec exploits the structure of a dataset to learn embeddings in a more sensible way than existing methods. This yields performance improvements in our ranking and prediction tasks. To the extent of our knowledge, Self-Supervised Feat2Vec is the first method to calculate continuous representations of data with multi-modal feature types and no explicit labels.

Future work could study how to reduce the amount of human knowledge our approach requires; for example by automatically grouping features into entities, or by automatically choosing a feature extraction function. These ideas can extend to our codebase that we make available⁶. Though further experimentation is necessary,

⁶The Feat2Vec model code is available here. The experi-

ments using non-proprietary data are available here.

References

- Bojanowski, P.; Grave, E.; Joulin, A.; and Mikolov, T. 2017. Enriching word vectors with subword information. *Transactions of the Association for Computational Linguistics* 5:135–146.
- Dyer, C. 2014. Notes on noise contrastive estimation and negative sampling. *arXiv preprint arXiv:1410.8251*.
- Dziugaite, G. K., and Roy, D. M. 2015. Neural network matrix factorization. *CoRR* abs/1511.06443.
- Guo, H.; TANG, R.; Ye, Y.; Li, Z.; and He, X. 2017. Deepfm: A factorization-machine based neural network for ctr prediction. In *Proceedings of the Twenty-Sixth International Joint Conference on Artificial Intelligence, IJCAI-17*, 1725–1731.
- Gutmann, M., and Hyvärinen, A. 2010. Noise-contrastive estimation: A new estimation principle for unnormalized statistical models. In *Proceedings of the Thirteenth International Conference on Artificial Intelligence and Statistics*, 297–304.
- Kalchbrenner, N.; Grefenstette, E.; and Blunsom, P. 2014. A convolutional neural network for modelling sentences. *CoRR* abs/1404.2188.
- Le, Q., and Mikolov, T. 2014. Distributed representations of sentences and documents. In *Proceedings of the 31st International Conference on Machine Learning (ICML-14)*, 1188–1196.
- Mikolov, T.; Chen, K.; Corrado, G.; and Dean, J. 2013a. Efficient estimation of word representations in vector space. *arXiv preprint arXiv:1301.3781*.
- Mikolov, T.; Sutskever, I.; Chen, K.; Corrado, G. S.; and Dean, J. 2013b. Distributed representations of words and phrases and their compositionality. In *Advances in neural information processing systems*, 3111–3119.
- Mnih, A., and Teh, Y. W. 2012. A fast and simple algorithm for training neural probabilistic language models. In *Proceedings of the International Conference on Machine Learning*.
- Rendle, S. 2010. Factorization machines. In *2010 IEEE 10th International Conference on Data Mining (ICDM)*, 995–1000. IEEE.
- Wu, L. Y.; Fisch, A.; Chopra, S.; Adams, K.; Bordes, A.; and Weston, J. 2018. Starspace: Embed all the things! In *Thirty-Second AAAI Conference on Artificial Intelligence*.
- Zhang, Y., and Wallace, B. 2017. A sensitivity analysis of (and practitioners’ guide to) convolutional neural networks for sentence classification. In *Proceedings of the Eighth International Joint Conference on Natural Language Processing*, volume 1, 253–263.
- Zheng, L.; Noroozi, V.; and Yu, P. S. Joint deep modeling of users and items using reviews for recommendation. In *Proceedings of the Tenth ACM International Conference on Web Search and Data Mining*.

ments using non-proprietary data are available here.